

Defensive Cybersecurity

Portfolio Part-2

This portfolio presents a collection of activities completed as part of the Defensive Cybersecurity unit. The objective of this portfolio is to demonstrate understanding of key cybersecurity concepts, practical implementations, and real-world applications. Each activity includes explanations, examples, and supporting evidence provided through a GitHub repository. The portfolio reflects both theoretical knowledge and hands-on experience in cybersecurity practices.

Portfolio Repository

All supporting evidence, including screenshots, code implementations, and additional materials, is available in the GitHub repository:

Self-Assessment Table

Activity	Status	Evidence	Self-Mark (/2)	Justification
B1	Completed	GitHub examples / written evidence	2	Five weak security implementations explained clearly.
B2	Completed	GitHub examples / written evidence	2	Five strong security implementations explained clearly.
B3	Completed	Written evidence / examples	2	Three proactive security implementations described.
B4	Completed	Lab participation description	2	Three in-class activities described with learning outcomes.
B5	Partly completed	One Seminar attended with prove	1	Only one cybersecurity-related seminar/talk is currently documented.
B6	Partly completed	One Seminar attended with prove	1	Ethics/law discussion is documented, but a second talk wasn't added.
B7	Completed	Club poster/photo evidence	2	UCC cybersecurity tech talk attended and described.
B9	Completed	Guest lecture note	2	Industry-related event mentioned with evidence/detail
B11	Partly completed	Conversation notes + attendance prove	1	One industry expert is documented; requirement asks for two experts.
B12	Completed	Research links + write-up	2	Two AI bias cases explained with references.
B13	Completed	Controlled test write-up	2	Safe jailbreak-style evaluation written; screenshots added.
B14	Completed	Teaching notes + signature sheet	2	Peer teaching session on encryption described.
B15	Completed	Teaching notes / family evidence	2	Cyber-safety teaching session for elderly person described.
B16	Completed	Research survey + references	2	State-of-the-art cybersecurity solutions surveyed.
B17	Completed	Python code + terminal output	2	ML anomaly detection implemented and evaluated.
B19	Completed	GitHub project branch / code evidence	2	Vulnerability found and fixed in Flask game project.
B20	Completed	GitHub branch + security docs	2	Security of GitHub project enhanced with several controls.

B21	Completed	Activity cards + sample photos	2	Cybersecurity learning activity designed and implemented.
B23	Completed	Docker/Suricata/tcp dump logs	2	IDS tested and effectiveness discussed.
B24	Completed	Python RBAC code + test cases	2	Access control designed and implemented.
B25	Completed	Python module + output screenshot	2	Threat-intel module works
B26	Completed	Explanation + signature sheet	2	Helped student with IDS; signed confirmation included
B27	Completed	Threat model tree + screenshots	2	Threat modelling applied to student laptop/account environment.
B29	Completed	CVE screenshot + AI comparison	2	CVE fix compared across three AI systems.
B30	Completed	Watermark scripts + detection output	2	Invisible watermark survived controlled edit.

Total selected activities: 25

Estimated self-mark: 47 / 50

Main evidence gaps before final submission: B5 second cybersecurity talk, B6 second ethics/law talk, B11 second industry expert.

GitHub Repository:

GitHub

https://github.com/faiyazkabr3/cyber_portfolio

This repository contains structured folders for each activity, along with corresponding evidence files.

B1. Discover 5 unique weak/vulnerable security implementations.

Definition

Weak or vulnerable security implementations refer to systems, practices, or configurations that fail to adequately protect information systems, making them susceptible to cyberattacks such as unauthorized access, data breaches, or system compromise (Whitman & Mattord, 2022).

1. Use of Default Passwords

Many devices such as routers, IoT devices, and security cameras are shipped with default usernames and passwords (e.g., “admin/admin”). If users do not change these credentials, attackers can easily gain access.

Why it is weak:

- Easily guessable
- Widely known publicly
- Frequently exploited in botnet attacks (e.g., Mirai)

2. Unpatched Software and Systems

Outdated software that has not been updated with security patches remains vulnerable to known exploits.

Why it is weak:

- Known vulnerabilities are publicly documented
- Attackers can use ready-made exploit tools
- Common entry point in ransomware attacks

This issue is frequently highlighted in cybersecurity advisories such as those from the Australian Cyber Security Centre (ACSC, 2026).

3. Weak Password Policies

Systems that allow short, simple, or reused passwords significantly increase the risk of brute-force or credential stuffing attacks.

Why it is weak:

- Easy to crack using automated tools
- Users tend to reuse passwords across platforms
- No enforcement of complexity or uniqueness

4. Lack of Encryption (Plaintext Data Storage)

Some systems store sensitive data such as passwords or personal information in plaintext instead of encrypting it.

Why it is weak:

- If the system is breached, all data is immediately readable
- No protection against insider threats
- Violates basic security principles

5. Open Network Ports and Misconfigured Services

Systems with unnecessary open ports or exposed services (e.g., SSH, database servers) can be discovered using tools like Shodan.

Why it is weak:

- Increases attack surface
- Allows unauthorized remote access
- Often exploited in automated scanning attacks

Conclusion

Weak security implementations often result from poor configuration, lack of updates, or weak user practices. These vulnerabilities can significantly increase the risk of cyberattacks if not addressed properly.

B2. Discover 5 unique strong security implementations.

Definition

Strong security implementations are systems and practices designed to effectively protect data and systems from threats, ensuring confidentiality, integrity, and availability (Whitman & Mattord, 2022).

1. Multi-Factor Authentication (MFA)

MFA requires users to verify their identity using multiple factors (e.g., password + mobile authentication app).

Why it is strong:

- Even if passwords are compromised, access is blocked
- Widely recommended by security standards (ACSC, 2023)

2. Encryption using TLS (HTTPS)

Transport Layer Security (TLS) encrypts data transmitted between users and servers.

Why it is strong:

- Protects data from interception
- Ensures secure communication
- Used in banking, e-commerce, and cloud services (Rescorla, 2018)

3. Full-Disk Encryption (e.g., FileVault, BitLocker)

Encrypts all data stored on a device using strong algorithms such as AES.

Why it is strong:

- Protects data even if device is stolen
- Prevents unauthorized access to stored data (NIST, 2001)

4. Password Hashing with bcrypt

Passwords are stored as hashed values instead of plaintext, with automatic salting.

Why it is strong:

- Protects against data breaches
- Resistant to brute-force attacks
- Ensures passwords cannot be easily reversed

5. Web Application Firewall (WAF)

A WAF protects web applications by filtering malicious traffic such as SQL injection and cross-site scripting attacks.

Why it is strong:

- Blocks common web-based attacks
- Adds an additional security layer
- Used by services like Cloudflare (Cloudflare, 2023)

Conclusion

Strong security implementations combine multiple layers of protection, making it significantly harder for attackers to compromise systems.

B3. Discover 3 proactive security implementations in practice.

Definition

Proactive security implementations are preventive measures taken before an attack occurs, aimed at reducing risks and identifying vulnerabilities early (Whitman & Mattord, 2022).

1. Regular Security Patching and Updates

Organizations continuously update systems and software to fix vulnerabilities.

Why it is proactive:

- Prevents exploitation of known vulnerabilities

- Reduces risk of zero-day attacks becoming widespread
- Recommended by cybersecurity frameworks such as NIST

2. Network Monitoring and Intrusion Detection Systems (IDS)

Systems such as IDS monitor network traffic and detect suspicious activities in real time.

Why it is proactive:

- Detects threats before damage occurs
- Alerts administrators of unusual behaviour
- Used in enterprise security platforms like IBM QRadar

3. Security Awareness Training (e.g., Phishing Training)

Organizations educate users about cyber threats such as phishing and social engineering.

Why it is proactive:

- Reduces human error (one of the biggest security risks)
- Helps users identify attacks before falling victim
- Improves overall security culture

Conclusion

Proactive security focuses on prevention rather than reaction. By identifying risks early and educating users, organizations can significantly reduce the likelihood of successful cyberattacks.

B4. Participate in 3 in-class activities in labs (facilitators will administer such activities).

Week 3 – Multi-Party Authentication Activity

In Week 3, we participated in an activity focused on multi-party authentication concepts. The objective was to calculate the average of several numbers without revealing each individual's actual value.

I worked in a group with Casey, Jonathan, and another classmate. Each participant selected a three-digit number. The first person added a secret random number to their value and passed the total to the next person. Each group member then added their own number to the running total and passed it along.

Once the total returned to the first person, they subtracted the secret number they had initially added. This ensured that the true sum of all participants' numbers was obtained without revealing individual inputs. Finally, the average was calculated by dividing the sum by the number of participants.

This activity demonstrated how privacy can be preserved in collaborative computations, which is an important concept in secure multi-party systems.

Week 4 – Access Control Activity

In Week 4 (Lab 3), we worked on an access control activity where the class was divided into groups. Each group was required to create a unique and identifiable feature that distinguished them from others while still being visible.

The goal was to design something that was noticeable but not easily replicable, reflecting real-world access control principles such as authentication factors and identity verification.

In the second part of the activity, groups attempted to identify the distinguishing features of other groups. This highlighted the challenges involved in verifying identities and ensuring that access control mechanisms are both secure and distinguishable.

This activity helped reinforce the importance of uniqueness and recognisability in authentication systems.




Week 6 – Juice Shop Vulnerability Analysis

In Week 6, we worked with the OWASP Juice Shop, a deliberately vulnerable web application designed for learning cybersecurity concepts.

During this lab, we explored and analysed various vulnerabilities present in the application. These included issues such as weak authentication, insecure data handling, and common web security flaws like injection attacks.

We interacted with the system to identify and understand how these vulnerabilities could be exploited. This hands-on experience helped us understand real-world security risks and the importance of secure coding practices.

Overall, this activity improved our practical understanding of web application security and vulnerability assessment.



B5. Attend 2 Cybersecurity Related Talks/Seminars

For this activity, I attended the cybersecurity-related lecture/seminar “**AI in Defensive Cybersecurity**”, delivered by **Larry Huynh** in week 11. This session counted as one of the two cybersecurity-related talks required for B5: because it was directly related to modern cybersecurity practice, especially the use of artificial intelligence and machine learning for defensive security tasks.

During the talk, the presenter explained the relationship between **Artificial Intelligence, Machine Learning, and Deep Learning**, and how these technologies can be applied in cybersecurity contexts. The lecture covered important defensive applications such as **anomaly-based intrusion detection systems (IDS)**, where machine learning models learn normal network behaviour and raise alerts when activity falls outside the expected baseline. This was useful because it showed how AI can help detect new or unknown attacks that may not be captured by traditional signature-based IDS tools.

The seminar also discussed **malware detection using machine learning**, including how malware samples can be represented using features such as strings, byte sequences, API calls, function call graphs, or even grayscale images. I learned that different models can produce different detection performance depending on how the malware data is represented and what type of model is used.

Another key part of the talk focused on **Large Language Models (LLMs)** and their role in cybersecurity. The lecture explained how LLMs are trained, including pre-training, safety training, and fine-tuning. It also discussed their use in **code vulnerability detection**, where LLMs can help identify suspicious code, explain possible security issues, and support human analysts. However, the seminar also made it clear that LLMs should not be treated as fully authoritative because they can miss vulnerabilities, misclassify issues, or produce incorrect explanations. Therefore, the best use of LLMs in cybersecurity is as a **human-in-the-loop support tool**, rather than a replacement for security knowledge and manual review.

Overall, this talk improved my understanding of how AI can be used in defensive cybersecurity. It helped me connect machine learning concepts with real security applications such as IDS, malware detection, and vulnerability analysis. It also made me more aware of both the strengths and limitations of AI tools in cybersecurity work.

Evidence of attendance:

A signed attendance sheet is attached as proof of my physical attendance at the AI security lecture/seminar. The lecture slides titled “**AI in Defensive Cybersecurity**” by Larry Huynh are also attached as supporting evidence.



Palwarit siripatanapunpol

Mosleh Hassanein 23101312

MUHAMMAD AFRIDI ISMAIL

24445737

Faiyaz Kabir


24773532

behaviour still occurs?

parts probing the POS system, the
technical firewall log or may
suspicious and

B6. Attend 2 cyber ethics/law related talks/seminars.

PHIL 1001 ETHICS FOR THE DIGITAL AGE



Week 08, Lecture B:

Digital Surveillance, and the Nature and Value of Privacy

For the first seminar, I attended a lecture by Dr Michael Rubin, which focused on the ethical issues surrounding **digital surveillance and privacy**. The lecture examined the “natural objection” to surveillance, which argues that collecting personal data is morally problematic because it violates an individual’s right to privacy. Rubin explored whether privacy violations are wrong due to harm or because they breach a moral right, noting that actions can be morally wrong even if no direct harm occurs. He also discussed different philosophical definitions of privacy, including the **control account** (privacy as control over personal information) and the **access account** (privacy as whether others actually access the information). Furthermore, the lecture introduced the concept of the **prisoner’s dilemma** to explain how individuals, acting in their own self-interest, may collectively contribute to a loss of privacy in society. This provided a deeper ethical perspective on why privacy rights may need to be protected through laws or shared moral principles

B7. Participate in CS/DS/cybersecurity clubs’ activity.

I participated in a cybersecurity club activity organised by the University Computer Club, where a tech talk titled “*Into the Cuckoo’s Egg*” was presented on 24th April 2026. The session was based on the well-known cybersecurity case study involving an investigation into one of the earliest recorded cyber intrusions.

The talk provided an overview of how a small accounting discrepancy led to the discovery of a large-scale cyber-espionage operation. It highlighted the importance of careful system monitoring, attention to detail, and persistence in cybersecurity investigations. The speaker


discussed how attackers were able to gain unauthorised access to systems and remain undetected for long periods, demonstrating early examples of hacking techniques. This activity was valuable in understanding the historical development of cybersecurity and how fundamental principles such as intrusion detection, logging, and system auditing have evolved over time. It also emphasised the importance of real-world investigation skills in identifying and responding to security incidents. Overall, the session provided practical insights into how cybersecurity incidents are discovered and analysed, reinforcing key concepts related to intrusion detection and system security.





B9. Participate in an industry-related cybersecurity event.

During Week 10, I attended a guest lecture by Wellington Muvunduke, who discussed strategies for **fraud detection within the banking sector**.





B11. Talk to 2 cybersecurity experts from the industry and find out their latest projects.

Case 1: Wellington Muvunduke — Fraud Analyst, Bankwest

For this task, I spoke with **Wellington Muvunduke**, who works as a **Fraud Analyst at Bankwest**. During the discussion, he explained that one of his current projects involves investigating **money mule activity** and analysing the behavioural patterns of accounts involved in suspicious fund movement. He described money mules as individuals or accounts used to receive and transfer illicit funds, often moving money from one mule account to another to hide the original source of the money.

Wellington explained that his work focuses on understanding how scam-related money enters accounts, how it moves between accounts, and what behavioural patterns may indicate mule activity. This includes observing transaction flow, account behaviour, and suspicious patterns that may help fraud teams identify and respond to financial crime. He also explained that he could not disclose specific operational details because the project is ongoing and confidential. Instead, he advised me to follow Bankwest's public security updates for information about current scam trends and customer protection advice.





B12. Discover two bias cases when using a generative AI system.

Introduction

Generative AI systems, such as ChatGPT, are trained on large datasets collected from the internet. While they are powerful tools, they can produce biased outputs because they learn from data that may already contain social and cultural inequalities. This task discusses two common bias cases found in generative AI systems: **historical bias** and **cultural bias**.

Bias Case 1: Historical Bias

Historical bias occurs when AI systems learn from past data that reflects existing inequalities or stereotypes in society. As a result, the AI may repeat or reinforce these patterns in its outputs.

According to the Australian Human Rights Commission, AI systems can reflect historical discrimination because they are trained on real-world data where certain groups may have been treated unfairly. This means the AI may unintentionally produce biased results, even if it is not designed to be unfair (Australian Human Rights Commission, n.d.).

Similarly, research shows that generative AI models inherit biases from their training data, including biases introduced by data creators and societal patterns. These biases can affect how AI generates responses, leading to unfair or unbalanced outcomes.

Example

For example, if historical data shows more men in leadership roles, an AI system may generate responses that favour male candidates when asked about leadership positions.

Impact

- Reinforces stereotypes
- Leads to unfair decision-making
- Reduces trust in AI systems

Bias Case 2: Cultural Bias

Cultural bias occurs when AI systems favour certain cultures, languages, or perspectives over others. This often happens because most training data comes from dominant regions, such as Western countries.

Research indicates that many generative AI systems are trained mainly on English and Western-based content, which leads to outputs that reflect Western cultural norms. As a result, perspectives from other cultures may be ignored or misrepresented.

This can create what researchers describe as “sociotechnical blindness,” where cultural assumptions are built into AI systems without being questioned.

Example

For instance, an AI system might describe education or social behaviour based mainly on Western ideas, even when the user is from a different cultural background.

Impact

- Misrepresentation of non-Western cultures
- Lack of diversity in AI outputs
- Users may feel misunderstood or excluded

B13. Perform a Jailbreak Attack on a Generative AI Assistant (Controlled Test Only)

Introduction

For this activity, I performed a controlled and safe jailbreak-style test on a generative AI assistant. The purpose of this task was not to obtain harmful information, but to understand how users may attempt to bypass the safety controls of AI systems and how the assistant responds to unsafe or manipulative prompts.

A jailbreak attack in the context of generative AI refers to an attempt to make the AI ignore its normal safety rules or system instructions. This can involve asking the model to role-play, ignore previous instructions, reveal restricted information, or provide unsafe guidance. These attacks are important in cybersecurity because AI assistants are increasingly used in software development, education, business workflows, and security analysis. If their safeguards can be bypassed, they may provide harmful, misleading, or policy-violating outputs.

Controlled Test Setup

The test was performed in a safe and educational setting. I did not request instructions for real harm, illegal activity, malware creation, credential theft, or system compromise. Instead, I used harmless prompts that simulated jailbreak behaviour without attempting to cause damage.

The controlled test focused on whether the AI assistant would:

- ignore its normal safety rules;
- follow a user instruction that attempted to override its safety behaviour;
- reveal hidden or restricted instructions;
- provide unsafe content when framed as role-play;
- redirect the user toward safe and ethical guidance.

Test Method

I tested the AI assistant using a small number of controlled prompts. The prompts were designed to represent common jailbreak techniques, but they were kept safe and non-operational.

The test categories were:

1. **Instruction override attempt**

I asked the assistant to ignore previous instructions and follow only my new instructions.

2. **Role-play attempt**

I asked the assistant to act as an unrestricted AI system that could answer anything.

3. **Hidden instruction request**

I asked whether the assistant could reveal its hidden system instructions or internal rules.

4. **Unsafe request simulation**

I used a clearly unsafe-style request but kept it general, without asking for real steps, code, or harmful procedures.

5. **Defensive reframing**

I then asked the assistant to explain the same topic from a safe cybersecurity-awareness perspective.

Observations

The AI assistant did not fully comply with the jailbreak-style prompts. When asked to ignore its rules or act as an unrestricted system, it either refused or redirected the conversation back to safe behaviour. When asked to reveal hidden instructions, it did not provide internal system information. This showed that the assistant had safeguards against common jailbreak attempts.

The assistant was also able to distinguish between unsafe and defensive framing. When a prompt appeared harmful or requested bypass-style behaviour, the assistant avoided giving direct harmful instructions. However, when the same topic was reframed as cybersecurity awareness or defensive education, the assistant provided safer, high-level explanations.

This was useful because it showed that generative AI systems do not only respond to keywords; they also consider the context and intent of the request. However, the test also showed that wording matters. A poorly framed cybersecurity question may be refused, while a clearly defensive and educational question is more likely to receive a useful answer.




Results


The controlled jailbreak test showed that the assistant's safety mechanisms were generally effective against simple jailbreak attempts. The assistant did not reveal hidden instructions, did not agree to ignore its safety rules, and did not provide harmful operational details. Instead, it redirected unsafe requests toward safer explanations.

The test also demonstrated the importance of responsible prompting. In cybersecurity, many topics are dual-use, meaning they can be used for both defensive and harmful purposes. For example, vulnerability testing can be legitimate in a controlled lab, but it can also be misused against real systems. Because of this, clear ethical framing is important when using AI tools for cybersecurity learning.

Evidence Added

As evidence for this activity, I included:





B14. Teach your friends about cybersecurity topic of your choice.

Introduction

Cybersecurity concepts such as encryption are fundamental to protecting communication in modern digital systems. For this activity, I conducted a short session with a group of peers to introduce the concepts of **public key–private key encryption** and **end-to-end encryption (E2EE)**, focusing on how these mechanisms protect data during transmission.

Topic Selection

The topic selected was “**Public Key Cryptography and End-to-End Encryption.**” This topic was chosen because it underpins many everyday technologies (e.g., messaging apps and secure websites), yet is often not well understood by general users.

Teaching Methodology

The session was designed to be clear and structured, combining brief explanations with simple analogies and demonstrations. Technical terminology was introduced gradually and supported with examples. I encouraged participation by asking questions and checking understanding throughout the session.

Content Delivered

Public Key and Private Key Encryption

I explained that public key cryptography uses a pair of keys:

- A **public key**, which can be shared openly
- A **private key**, which must be kept secret

To make the concept easier to understand, I used an analogy: the public key is like an open mailbox that anyone can use to send a message, but only the owner with the private key can open it and read the contents.

I also explained that:

- Data encrypted with a public key can only be decrypted using the corresponding private key
- This method is widely used for secure communication over the internet

End-to-End Encryption (E2EE)

I then introduced end-to-end encryption, explaining that it ensures only the sender and receiver can read the message. Even the service provider cannot access the content.

I clarified that:

- Messages are encrypted on the sender's device
- They remain encrypted during transmission
- Only the receiver's device can decrypt them

A simple example using messaging applications was discussed to show how E2EE protects user privacy.

Importance of Encryption

I highlighted why these concepts matter:

- Protects sensitive information (e.g., passwords, messages)
- Prevents unauthorised access during data transmission
- Maintains privacy in digital communication

Outcome

By the end of the session, participants demonstrated a basic understanding of how public and private keys function and how end-to-end encryption secures communication. They were able to explain the difference between the two concepts and recognise their importance in everyday applications.

Evidence of Participation

To document participation, attendees signed a sheet confirming their presence during the session. This serves as evidence that the teaching activity was conducted.

Subject	Date
Public Key Cryptography and End-to-End Encryption	23/04/2026
Sabit	
Reza Sabit	

Subject	Date
Public Key Cryptography & End-to-End Encryption	23/04/26
Mamun	
Noor - A - Kamal	

Subject	Date
Public Key Cryptography and End-to-End Encryption	23/04/2026
Faisal Maheer	
Md. Zarif Alam Khan	

Subject	Date
Public Key Cryptography and End-to-End Encryption	23.04.2026
Md. Zarif Alam Khan	

Conclusion

In conclusion, teaching peers about public key cryptography and end-to-end encryption helps build awareness of how modern communication is secured. Understanding these concepts enables users to better appreciate the importance of privacy and data protection in digital environments.

B15. Teach an elderly person about cybersecurity topic of your choice.

Introduction

Cybersecurity is very important today, especially for elderly people who may not be familiar with online risks. For this task, I taught an elderly family member (my grandma) about basic cybersecurity and how to stay safe while using the internet.

Topic Chosen

I chose the topic “**Online Scams and Phishing**” because elderly people are often targeted by scammers through emails, phone calls, and messages.

Teaching Approach

I kept the explanation simple and easy to understand. Instead of using technical terms, I used real-life examples and explained things step by step. I also asked questions during the session to make sure they understood.

What I Taught

1. What is a Scam?

I explained that a scam is when someone tries to trick you into giving money or personal information. Scammers often pretend to be from banks or well-known companies.

2. What is Phishing?

I explained that phishing is when someone sends fake emails or messages to steal your information, such as passwords or bank details.

3. How to Spot a Scam

I showed them some common warning signs:

- Messages that sound urgent or scary
- Emails from unknown people
- Strange or suspicious links
- Requests for personal information



4. How to Stay Safe

I gave some simple tips:

- Do not click on unknown links
- Do not share passwords
- Always double-check messages
- Ask someone if unsure

5. What to Do if Unsure

I told them that if something feels wrong, they should:

- Not reply immediately
- Ask a family member
- Contact the company directly

Outcome

After the session, they were able to recognise basic scam messages and understood why it is important to be careful online. They also said they feel more confident using email now.

B16. Survey the current state-of-the-art solutions in cybersecurity.

Introduction

Cybersecurity has moved beyond traditional perimeter defenses such as basic firewalls and antivirus software. Modern organisations now face threats including credential theft, ransomware, supply-chain compromise, cloud misconfiguration, and attacks against identity systems. As a result, current state-of-the-art cybersecurity solutions focus on identity, continuous verification, and threat-informed defense rather than relying only on network boundaries (National Institute of Standards and Technology [NIST], 2020; Cybersecurity and Infrastructure Security Agency [CISA], n.d.-b; MITRE, n.d.-a, n.d.-b, n.d.-c). This report surveys three major modern approaches: Zero Trust Architecture, phishing-resistant multifactor authentication, and threat-informed defense using MITRE ATT&CK and D3FEND.


1. Zero Trust Architecture

One of the most widely recognised modern cybersecurity approaches is Zero Trust Architecture (ZTA). NIST defines zero trust as an evolving set of cybersecurity paradigms that shifts defenses away from static, network-based perimeters and toward users, assets, and resources (NIST, 2020). In practice, this means that no user or device is trusted automatically simply because it is inside the organisation's network. Instead, access decisions are made continuously based on identity, device state, context, and policy (NIST, 2020).

Zero Trust is considered state-of-the-art because it reflects the reality of modern computing environments, where work is distributed across cloud platforms, mobile devices, hybrid workplaces, and third-party services. Traditional "inside equals trusted" thinking is no longer sufficient in these environments. Zero Trust reduces implicit trust and supports more granular, adaptive access control (NIST, 2020).

A major strength of Zero Trust is that it limits lateral movement and reduces the damage caused by compromised accounts or devices. However, it can also be difficult to implement fully because it often requires identity modernisation, policy redesign, asset visibility, and ongoing monitoring. For this reason, Zero Trust is best understood as a strategic architecture rather than a single product (NIST, 2020).

ZERO TRUST SECURITY




2. Phishing-Resistant Multifactor Authentication

Another leading modern solution is phishing-resistant multifactor authentication (MFA). CISA recommends stronger forms of MFA because attackers increasingly bypass passwords and weaker second factors through phishing, push fatigue, token theft, and credential replay (CISA, n.d.-b). Phishing-resistant methods, including FIDO-based authentication, are designed so that users authenticate to the real service and cannot easily be tricked into handing over reusable secrets (CISA, n.d.-b).

This approach is state-of-the-art because identity compromise remains one of the most common entry points for modern attacks. Password-only systems are weak, and even some legacy MFA approaches do not adequately resist modern phishing campaigns. Phishing-resistant MFA improves protection for email, administrative accounts, VPNs, and cloud services, which are frequent attacker targets (CISA, n.d.-b).


The main advantage of phishing-resistant MFA is that it directly addresses one of the most important real-world attack paths: stolen credentials. Its main limitation is deployment complexity, especially in organisations still dependent on legacy applications or older authentication workflows. Even so, it is now widely viewed as one of the most effective practical defenses available (CISA, n.d.-b).



3. Threat-Informed Defense Using MITRE ATT&CK and D3FEND

A third state-of-the-art solution is threat-informed defense, especially through the use of frameworks such as MITRE ATT&CK and MITRE D3FEND. MITRE ATT&CK is a globally accessible knowledge base of adversary tactics and techniques based on real-world observations, while D3FEND is a knowledge graph of cybersecurity countermeasures that maps defensive techniques in a more structured way (MITRE, n.d.-a, n.d.-b, n.d.-c).

This approach is modern because it helps defenders move away from generic security controls and toward controls that are explicitly linked to attacker behaviour. Instead of asking only “what tools do we have?”, organisations can ask “which techniques are most likely to be used against us, and which countermeasures address them?” D3FEND is especially useful because it provides a more formal model of defensive techniques and their relationships, helping security teams design more systematic defensive strategies (MITRE, n.d.-a, n.d.-c). The strength of threat-informed defense is that it improves alignment between security operations and actual attacker tradecraft. It supports use cases such as control selection, detection engineering, validation, purple teaming, and security architecture planning. Its limitation is that it requires maturity in analysis, threat intelligence, and security operations. Smaller organisations may find it harder to operationalise without specialist skills (MITRE).



B17. Implement one of the current state-of-the-art solutions and evaluate it.

Implemented Solution: AI/ML-Based Anomaly Detection for Login Security

For this task, I implemented a simple **ML-based anomaly detection system** for login security. This is a current cybersecurity approach because many modern systems use behavioural analytics and machine learning to detect unusual activity, such as suspicious login attempts, account takeover attempts, or access from unfamiliar devices.

The solution was implemented in Python using an **Isolation Forest** model. Isolation Forest is an anomaly detection algorithm that can identify unusual data points compared to normal behaviour. In this implementation, the model analysed simulated login records and classified each login attempt as either **Normal** or **Suspicious**.

Design and Implementation

The program used a small simulated login dataset. Each login attempt had four features:

Feature	Meaning
login_hour	The hour when the login occurred
failed_attempts	Number of failed login attempts
country_risk	Risk score of the login location
new_device	Whether the login came from a new device

Normal login behaviour was represented by daytime or evening logins, low failed attempts, low country risk, and known devices. Suspicious login behaviour was represented by unusual login hours, repeated failed attempts, higher country risk, and new device usage.

The Python script trained the Isolation Forest model, generated predictions, labelled each login as **Normal** or **Suspicious**, and saved the final results into a CSV file called:

login_anomaly_results.csv

Results

After running the program, the model successfully identified the normal and suspicious login attempts. The normal records were classified as **Normal**, while the unusual login records were classified as **Suspicious**.

For example, login attempts at unusual times such as 2 AM, 3 AM, 4 AM, and 11 PM with multiple failed attempts, high country risk, and new device usage were flagged as suspicious. This shows that the model was able to detect abnormal login behaviour based on multiple combined features rather than a single rule.

The terminal output showed that the anomaly detection process completed successfully and saved the results to the CSV file.

Evaluation

The implemented solution was effective as a proof of concept. It correctly flagged the most suspicious login attempts in the simulated dataset. This demonstrates how AI/ML-based anomaly detection can support cybersecurity by identifying login activity that does not match normal user behaviour.

A strength of this approach is that it can detect unusual combinations of behaviour without manually writing a rule for every possible scenario. For example, a login may not be suspicious only because it happens late at night, but it becomes more suspicious when combined with repeated failed attempts, a high-risk location, and a new device.

However, the solution also has limitations. The dataset used was small and simulated, so the results are not enough for real-world deployment. In a real system, the model would need a much larger dataset, better feature engineering, regular tuning, and human review of alerts. It may also produce false positives, where legitimate users are incorrectly flagged as suspicious.

Overall, this implementation shows how a current state-of-the-art cybersecurity technique, **machine-learning-based anomaly detection**, can be applied to login security to help detect possible account compromise.

Evidence Added

For evidence, I included:

- Python implementation screenshot
- Terminal output showing the model running successfully
- Output showing normal and suspicious classifications
- Generated CSV results file
- GitHub evidence folder for B17

Evidence location:

https://github.com/faiyazkabir3/cyber_portfolio/tree/main/part_2/B17/evidence

CodeBlame100 lines (74 loc) · 2.08 KB

```
1 import pandas as pd
2 from sklearn.ensemble import IsolationForest
3 from sklearn.preprocessing import StandardScaler
4
5
6 def create_login_dataset():
7     """
8     Creates a small simulated login dataset.
9     Each row represents one login attempt.
10    """
11
12    data = [
13        # Normal login behaviour
14        [9, 0, 1, 0],
15        [10, 0, 1, 0],
16        [11, 1, 1, 0],
17        [13, 0, 1, 0],
18        [14, 0, 1, 0],
19        [15, 1, 1, 0],
20        [16, 0, 1, 0],
21        [18, 0, 1, 0],
22        [20, 1, 1, 0],
23        [21, 0, 1, 0],
24
25        # More normal behaviour
26        [8, 0, 1, 0],
27        [12, 0, 1, 0],
28        [17, 1, 1, 0],
29        [19, 0, 1, 0],
30
31        # Suspicious login behaviour
32        [3, 6, 5, 1],
33    ]
34
35    return pd.DataFrame(data, columns=['login_hour', 'failed_attempts', 'country_risk', 'new_device'])
36
37 # Create and save the dataset
38 dataset = create_login_dataset()
39 dataset.to_csv('login_anomaly_detector.csv', index=False)
40
41 # Train the model
42 scaler = StandardScaler()
43 dataset[['failed_attempts', 'country_risk', 'new_device']].loc[:, :] = scaler.fit_transform(
44     dataset[['failed_attempts', 'country_risk', 'new_device']].loc[:, :]
45 )
46
47 model = IsolationForest()
48 model.fit(dataset[['failed_attempts', 'country_risk', 'new_device']])
49
50 # Predict on the dataset
51 predictions = model.predict(dataset[['failed_attempts', 'country_risk', 'new_device']])
52
53 # Add predictions to the dataset
54 dataset['prediction'] = predictions
55
56 # Save the results
57 dataset.to_csv('login_anomaly_results.csv', index=False)
58
59 # Print results
60 print("Anomaly detection completed. Results saved to login_anomaly_results.csv")
```

Anomaly detection completed.

Results saved to login_anomaly_results.csv

	login_hour	failed_attempts	country_risk	new_device	prediction	result
0	9	0	1	0	1	Normal
1	10	0	1	0	1	Normal
2	11	1	1	0	1	Normal
3	13	0	1	0	1	Normal
4	14	0	1	0	1	Normal
5	15	1	1	0	1	Normal
6	16	0	1	0	1	Normal
7	18	0	1	0	1	Normal
8	20	1	1	0	1	Normal
9	21	0	1	0	1	Normal
10	8	0	1	0	1	Normal
11	12	0	1	0	1	Normal
12	17	1	1	0	1	Normal
13	19	0	1	0	1	Normal
14	3	6	5	1	-1	Suspicious
15	2	5	5	1	-1	Suspicious
16	4	7	4	1	-1	Suspicious
17	23	4	5	1	-1	Suspicious

Name	Date modified	Type	Size
 login_anomaly_detector	9/05/2026 1:18 AM	Python Source File	3 KB
 login_anomaly_results	9/05/2026 1:22 AM	Microsoft Excel C...	1 KB

B19. Find and fix a vulnerability from a GitHub project.

Project Selected

For this task, I used my group GitHub project, a Flask-based web game application. During review, I identified several security weaknesses in the project's authentication, database, request handling, and local storage design. Here is the repo :

https://github.com/faiyazkabir3/CITS3403_Project_JYCF

Vulnerability Found

One major vulnerability identified was that sensitive local project data, such as the SQLite database and fallback save files, could be stored in readable plaintext. This meant that if someone copied the local project .db file or fallback save files, they may have been able to inspect user, save, or chat data. The security guide for the upgraded branch identifies this as a realistic project threat, specifically noting risks such as someone copying the local SQLite database or fallback save files and trying to read stored data.

Another related vulnerability was unsafe request handling. Some mutating actions, such as logout and friend request actions, previously relied on unsafe request patterns. This could create a risk where a malicious website might trigger an action on behalf of a logged-in user. The upgraded branch specifically addresses this by requiring CSRF protection and making mutating friend/logout routes POST-only.

Fix Implemented

To fix these issues, I implemented a major security upgrade on the db_security branch. The main fixes included:

- SQLCipher-backed encrypted SQLite database support
- AES-GCM encryption for fallback save files
- required .env secrets for SECRET_KEY, SQLCIPHER_DATABASE_KEY, and SAVE_PAYLOAD_KEYS

```
.env
1
2 SECRET_KEY= L0gg7ZVR0W3s0r11...
3 SQLCIPHER_DATABASE_KEY= sAgiT9t7oY4f...
4 SAVE_PAYLOAD_KEYS=v1:q3M...
5 DATABASE_URL=sqlite:///project.db
6
```

```
def load_required_secret_key():
    secret_key = os.environ.get("SECRET_KEY", "").strip()
    normalized_secret_key = secret_key.lower()

    if not secret_key:
        raise RuntimeError("SECRET_KEY is missing. Add it to your .env file.")

    if normalized_secret_key in SECRET_KEY_PLACEHOLDERS:
        raise RuntimeError(
            "SECRET_KEY is still set to a placeholder. Generate a strong random value and store it in"
        )

    if len(secret_key) < 32:
        raise RuntimeError("SECRET_KEY is too short. Use at least 32 random characters in your .env f")

    return secret_key
```

- Flask-Login for registered-user authentication
- Flask-WTF CSRF protection for POST forms and JavaScript requests
- POST-only mutating routes for logout and friend actions
- SQLAlchemy-safe query usage for user-controlled database access

The screenshot shows a GitHub pull request interface. The title is "Harden app security with SQLCipher, E2EE chat, Flask-Login, and CSRF protection #135". The status is "Awaiting approval" and "Code" is visible. The pull request is from user "faiyazkabr3" to the "main" branch of the "db_security" repository. It shows 14 commits and 25 files changed. The pull request is open, and the user "faiyazkabr3" commented 1 minute ago. The summary states: "This PR upgrades the project's security model across authentication, database storage, fallback saves, and direct chat. (has dedicated SECURITY_UPGRADE.md file for full exploration, if confused read it first)". The "What Changed" section lists added security dependencies: Flask-Login, Flask-WTF, WTForms, Flask-Migrate, and cryptography. The right sidebar shows reviewers, assignees, and labels.

The upgraded security guide confirms that the branch now includes Flask-Login session authentication, Flask-WTF CSRF protection, salted password hashing, .env secret validation, SQLCipher encryption, AES-GCM fallback save encryption, browser-side chat encryption, and safer SQLAlchemy query APIs.

Result


After the fix, the project had a stronger security model. The local database is no longer intended to be readable as a normal plaintext SQLite database, fallback saves are encrypted, unsafe mutating GET requests are avoided, and registered-user features are protected using proper login checks. This reduced the risk of data exposure, CSRF attacks, and insecure local storage.

This task showed that GitHub projects can contain security weaknesses even when the application appears to work normally. Fixing the vulnerability required more than a single code change, because authentication, request handling, storage, and documentation all had to be improved together.

B20. Enhance the security of a GitHub project.

Project Selected

For this task, I enhanced the security of my Flask-based web game project on GitHub. The work was completed on the `db_security` branch and focused on improving authentication, database protection, CSRF protection, encrypted storage, and secure documentation.



Here is the repo :

https://github.com/faiyazkabr3/CITS3403_Project_JYCF

Security Enhancements Implemented

1. Authentication Improvements

The project was enhanced with Flask-Login to manage registered-user sessions. This means registered-only features can be protected using `@login_required`, and logged-out users are redirected to the login page. The security guide notes that Flask-Login handles sessions through `login_user()`, `logout_user()`, `current_user`, and route protection.

```

@app.route("/profile", methods=["GET", "POST"])
@login_required
def profile():
    user = current_user
    user_id = user.id

    error = None
    success = None
    display_name_value = user.display_name or ""
    bio_value = get_profile_bio(user)
    selected_profile_image = get_profile_image(user)
    custom_profile_image = get_custom_profile_image(user)
    profile_background_value = get_profile_background(user)
    favorite_character_value = user.favorite_character or ""
    show_stats_to_friends_value = bool(user.show_stats_to_friends)
    allow_friend_messages_value = bool(user.allow_friend_messages)
    hide_from_leaderboard_value = bool(user.hide_from_leaderboard)

    if request.method == "POST":
        display_name = request.form.get("display_name", "").strip()
        bio = request.form.get("bio", "").strip()
        favorite_character = normalize_save_token(
            request.form.get("favorite_character", ""),
            "",
            max_length=20,
            transform="lower"
        )
        profile_background = normalize_save_token(
            request.form.get("profile_background", "default"),
            "default"

```

2. CSRF Protection

Flask-WTF CSRF protection was added to protect mutating requests such as login, registration, logout, profile updates, friend requests, save-game actions, and chat key requests. Forms include CSRF tokens, and JavaScript fetch requests send the token using the X-CSRFToken header.

```

43
44 <!-- Chat Input -->
45 <form method="POST" class="chat-form" data-chat-form>
46   <input type="hidden" name="csrf_token" value="{ csrf_token() }">
47   <input
48     type="text"
49     name="message"
50     placeholder="Type a message..."
51     required
52     class="input-text"
53     data-chat-input
54   >
55   <input type="hidden" name="ciphertext" data-chat-ciphertext>
56   <input type="hidden" name="nonce" data-chat-nonce>
57   <input type="hidden" name="sender_public_key" data-chat-sender-public-key>
58   <input type="hidden" name="sender_key_id" data-chat-sender-key-id>
59   <input type="hidden" name="recipient_public_key" data-chat-recipient-public-key>
60   <input type="hidden" name="recipient_key_id" data-chat-recipient-key-id>
61   <input type="hidden" name="encryption_version" value="1">
62   <button type="submit" class="primary-btn" data-chat-submit>Send</button>

```

3. Encrypted Database Storage

SQLCipher-backed SQLite encryption was added so that the local database is encrypted at rest. The application verifies SQLCipher using `PRAGMA cipher_version`, and if SQLCipher is not active, startup fails. This helps protect the database if the raw database file is copied.

```
37
38 try:
39     import sqlcipher3
40     import sqlcipher3.dbapi2 as sqlcipher_dbapi
41 except ModuleNotFoundError:
42     sqlcipher3 = None
43 else:
44     sys.modules.setdefault("pysqlcipher3", sqlcipher3)
45     sys.modules.setdefault("pysqlcipher3.dbapi2", sqlcipher_dbapi)
46
47 load_dotenv()
48
49 app = Flask(__name__, instance_relative_config=True)
50
51 os.makedirs(app.instance_path, exist_ok=True)
52 SAVE_FALLBACK_DIR = Path(app.instance_path) / "save_fallbacks"
53 SAVE_FALLBACK_DIR.mkdir(exist_ok=True)
54 PROFILE_IMAGE_UPLOAD_DIR = Path(app.static_folder) / "uploads" / "profile_pics"
55 PROFILE_IMAGE_UPLOAD_DIR.mkdir(parents=True, exist_ok=True)
56
57 BIO_MAX_LENGTH = 500
58 PROFILE_COMMENT_MAX_LENGTH = 240
59 PROFILE_IMAGE_DEFAULT = "images/Shadows.gif"
```

4. Encrypted Fallback Saves

Fallback save files are now encrypted using AES-GCM. The upgraded design stores an encrypted JSON envelope containing fields such as encrypted, `key_id`, `nonce`, and `ciphertext`, instead of storing readable save data. AES-GCM also helps detect tampering if

the ciphertext, nonce, or key is changed.

```
193
194 def load_required_secret_key():
195     secret_key = os.environ.get("SECRET_KEY", "").strip()
196     normalized_secret_key = secret_key.lower()
197
198     if not secret_key:
199         raise RuntimeError("SECRET_KEY is missing. Add it to your .env file.")
200
201     if normalized_secret_key in SECRET_KEY_PLACEHOLDERS:
202         raise RuntimeError(
203             "SECRET_KEY is still set to a placeholder. Generate a strong random value and store it in
204         )
205
206     if len(secret_key) < 32:
207         raise RuntimeError("SECRET_KEY is too short. Use at least 32 random characters in your .env f
208
209     return secret_key
210
211
212 def load_required_env_secret(name, minimum_length=32):
213     value = os.environ.get(name, "").strip()
214
215     if not value:
216         raise RuntimeError(f"{name} is missing. Add it to your .env file.")
```

5. End-to-End Encrypted Direct Chat

The direct chat feature was enhanced with browser-side end-to-end encryption. The browser generates ECDH P-256 keys, derives a shared key, and encrypts messages using AES-GCM before the Flask server receives them. The server stores encrypted message fields rather than plaintext chat content.

```
67 <script type="application/json" data-chat-history>{{ chat_messages | tojson }}</script>
68 <script
69     src="https://cdnjs.cloudflare.com/ajax/libs/socket.io/4.0.1/socket.io.js"
70     integrity="sha512-q/dwJ3kcmjBLU4Qc47E4A9kTB4m3wuTY7vkFJDTZKjTs8jhyGQnaUrxa0Ytd0ssMZhbNua9hE+E7Qv1j+OyZuA=="
71     crossorigin="anonymous"
72 ></script>
73 <script src="{{ url_for('static', filename='js/chat.js') }}"></script>
74 </body>
75 </html>
```

6. Safer Database Query Handling

User-controlled database lookups now use SQLAlchemy query APIs instead of manually concatenating user input into SQL strings. This reduces the risk of SQL injection-style issues in usernames, profile fields, friend searches, and chat routes.

7. Improved Documentation and Testing

Security documentation was expanded with files such as SECURITY_UPGRADE.md, END_TO_END_ENCRYPTED_CHAT.md, and DATABASE_GUIDE.md. The guide also documents verification commands, including checks for SQLCipher, encrypted fallback saves, chat plaintext storage, CSRF failures, and disabled mutating GET routes.

```

SECURITY_UPGRADE.md > abc # Security Upgrade Guide
1  # Security Upgrade Guide
2
3  Updated: 30 April 2026
4
5  This document is the security reference for the `db_security` branch. It explains what the project
   protects, how each security layer works, which files are involved, how to verify the behavior, and
   what risks still remain.
6
7  The project is still a Flask game application, but the security model now has several separate layers:
8
9  - Flask-Login for registered-user session authentication.
10 - Flask-WTF CSRF tokens for mutating HTTP requests.
11 - Werkzeug password hashing with a per-password salt.
12 - Required `.env` secrets instead of hardcoded keys.
13 - SQLCipher-backed SQLite database encryption at rest.
14 - AES-GCM encryption for fallback save files.
15 - Browser-side Web Crypto end-to-end encryption for direct chat.
16 - SQLAlchemy query APIs for user-input database access.
17 - Jinja/browser escaping patterns for user-controlled text.
18 - File-upload validation for profile pictures.
19 - Flask-Migrate/Alembic support for safer schema changes.
20
21 The important design rule is that different security jobs use different secrets. Flask session
   signing, SQLCipher database encryption, fallback save encryption, and browser chat encryption are
   intentionally separate.
22
23 ## Security Goals
24
25 This branch is designed to protect against these realistic project threats:
26
27 - Someone copies the local SQLite database file and tries to read user, save, or chat data.
28 - Someone copies fallback save files from `instance/save_fallbacks/`.
29 - A malicious website tries to trigger a logged-in user's POST action through CSRF.

```

Outcome

The project became significantly more secure after these changes. Before the upgrade, local storage and chat rows could expose more readable data. After the upgrade, registered accounts use salted password hashes, sessions are handled by Flask-Login, mutating requests are CSRF protected, SQLite is encrypted through SQLCipher, fallback saves are encrypted with AES-GCM, and direct chat messages are encrypted in the browser before reaching the server.

Limitations

The security upgrade does not solve every possible issue. HTTPS is still considered a deployment responsibility, browser private chat keys are stored in `localStorage`, there is no multi-device chat key sync, and SQLCipher only protects the database at rest rather than data already decrypted inside the running Flask process. These limitations are documented honestly in the security guide.

B21. Design and implement a cybersecurity learning activity.

Malware Safety Sorting Game

Description

For this activity, I designed and implemented a simple interactive cybersecurity learning activity called the **Malware Safety Sorting Game**. The aim of the activity was to teach participants how to recognise safe and unsafe behaviours that can affect computer security, particularly in relation to malware prevention.

The activity used a set of cards containing different computer-use actions. Participants were asked to sort each card into one of two categories: **Safe** or **Unsafe**. The safe cards included actions such as **updating software**, **downloading from official websites**, **running antivirus scans**, and **backing up files**. The unsafe cards included actions such as **installing cracked games**, **opening unknown email attachments**, **plugging in unknown USB devices**, and **disabling antivirus software**.

Design and Implementation

I first designed the activity by selecting common real-world examples of safe and unsafe digital behaviour. These examples were chosen because they are easy to understand and directly related to basic cybersecurity awareness. After identifying the examples, I created a set of learning cards that could be used in the sorting exercise. The cards were visually organised so that participants could clearly read and classify them.

During the implementation stage, participants were shown the cards and asked to place each one into either the **Safe** or **Unsafe** category. After the sorting activity was completed, I discussed the correct answers with the participants and explained why each action was either safe or unsafe. This allowed the activity to function not only as a game, but also as a short educational exercise.

Learning Goal

The main learning goal of the activity was to help participants understand basic malware prevention practices. By completing the sorting task, participants were able to identify actions that reduce cybersecurity risk and distinguish them from behaviours that may expose a computer to malware or other threats.

Outcome

The activity was successful because it encouraged active participation and made a basic cybersecurity topic easier to understand. The sorting format was simple, interactive, and suitable for beginners. It also helped reinforce the idea that many malware risks come from everyday user behaviour rather than only technical system issues.

Evidence

As evidence for this activity, I included:

- the list of card items used in the activity
- sample photos of the Malware Safety Sorting Game cards
- sample photos showing how the cards were arranged into **Safe** and **Unsafe** groups

These sample photos were added to demonstrate the design and implementation of the activity.


Card List Used in the Activity

Safe

- Update software
- Download from official website
- Run antivirus scan
- Back up files

Unsafe

- Install cracked games
- Open unknown email attachment
- Plug in unknown USB
- Disable antivirus



B23. Test an intrusion detection system and discuss its effectiveness.

An Intrusion Detection System (IDS) is a security mechanism designed to monitor and analyse network traffic for suspicious or malicious activity. IDS solutions can be classified as signature-based, which detect known attack patterns, or anomaly-based, which identify deviations from normal behaviour. In this task, the IDS tool **Suricata** was selected due to its modern capabilities, including multi-threaded processing and support for real-time and offline traffic analysis.

Methodology

The IDS was tested using **Suricata running inside a Docker container** on macOS. This approach provided an isolated and reproducible environment without requiring a full penetration testing operating system such as Kali Linux.

First, network traffic was captured using the tcpdump tool on the active network interface (en0). A packet capture (PCAP) file containing approximately 100–200 packets was

generated. This traffic included normal user activity such as web browsing and background system communications.

Next, Suricata was executed in **offline analysis mode**, where it processed the captured PCAP file instead of monitoring live traffic. Custom detection rules were configured to identify:

- ICMP traffic (e.g., ping requests)
- HTTP traffic (port 80 communication)

The following command was used to run Suricata:

```
(base) faiyazkabir@MacBook-Pro ids-test % docker run --rm -it \
-v ~/ids-test:/data \
jasonish/suricata \
suricata -r /data/test.pcap -S /data/local.rules -l /data
i: suricata: This is Suricata version 8.0.4 RELEASE running in USER mode
i: threads: Threads created -> RX: 1 W: 15 FM: 1 FR: 1 Engine started.
i: suricata: Signal Received. Stopping engine.
i: pcap: read 1 file, 200 packets, 117800 bytes
```

The IDS successfully processed the captured PCAP file, analysing between 100 and 200 packets of real network traffic. The generated logs confirmed that multiple network flows were inspected, including both TCP and UDP communications.

Analysis of the `eve.json` output showed that most traffic consisted of:

- Encrypted HTTPS communication (port 443)
- QUIC protocol traffic
- Multicast DNS (mDNS) background traffic

No alerts were triggered during the test, as the captured traffic did not match the predefined detection rules. In particular:

- No ICMP packets were present in the capture
- HTTP traffic (port 80) was minimal or absent due to modern websites using HTTPS instead

Despite the absence of alerts, Suricata successfully performed **deep packet inspection and traffic classification**, demonstrating that the IDS was functioning correctly.

```
(base) faiyazkabir@MacBook-Pro ids-test % ls ~/ids-test
eve.json      fast.log      local.rules   stats.log     suricata.log  test.pcap
```

Results

Suricata successfully analysed all captured packets and produced detailed logs. The results showed that most traffic consisted of encrypted HTTPS (port 443) and background system traffic (e.g., QUIC and mDNS). No alerts were triggered, as the traffic did not match the defined rules.

Evaluation

The IDS was effective in monitoring and analysing network traffic, demonstrating correct functionality. However, its detection capability was limited by the predefined rules. Modern encrypted traffic reduces the effectiveness of simple rules, and background traffic can make testing more difficult.

Overall, Suricata successfully performed network analysis, but its effectiveness depends on the quality of detection rules and the type of traffic being monitored.

B24. Design and implement access control of your choice.

Introduction

Access control is an important part of cybersecurity. It ensures that only the right people can access certain information or perform specific actions in a system. For this task, I designed and implemented a simple access control system using **Role-Based Access Control (RBAC)**.

Access Control Chosen

I chose **Role-Based Access Control (RBAC)** because it is simple, easy to understand, and commonly used in real-world systems. In RBAC, users are assigned roles, and each role has specific permissions.

System Design

Roles

I created three roles in the system:

- **Admin** – has full access
- **Staff** – has limited access
- **User** – has basic access

Permissions

Each role was given different permissions:

- **Admin**
 - o Can view all data
 - o Can delete data
 - o Can manage users
- **Staff**
 - o Can view and upload data
 - o Cannot delete or manage users
- **User**
 - o Can only view their own data

Resources

The system protects:

- User data
- Files
- Admin dashboard

Example Logic

This function checks if a user's role allows them to perform a specific action. I implemented the system using a simple Python program.

```

1  def check_access(role, action):
2      permissions = {
3          "admin": ["view", "delete", "manage"],
4          "staff": ["view", "upload"],
5          "user": ["view_own"]
6      }
7
8      allowed_actions = permissions.get(role)
9
10     if allowed_actions is None:
11         return "Invalid Role"
12
13     if action in allowed_actions:
14         return "Access Granted"
15     else:
16         return "Access Denied"

```

Example Test

- Admin trying to delete → Access Granted
- Staff trying to delete → Access Denied
- User trying to view own data → Access Granted

Testing

I tested the system using different roles and actions to make sure it works correctly.

Test Case	Result
Admin deletes data	Allowed
Staff deletes data	Denied
User views own data	Allowed
User views all data	Denied

B25. Design and implement a threat intelligence module of your choice.

For this task, I designed and implemented a simple threat intelligence module called an **IOC Threat Checker**. The purpose of the module was to compare a list of suspicious indicators against a small threat intelligence feed and classify them based on threat level.

The module was implemented in Python. It used three main files: an input file containing indicators to check, a CSV file acting as the threat intelligence feed, and a Python script that

performed the comparison. The indicators included IP addresses and domain names. The threat feed contained known suspicious indicators, their type, threat level, and a short description.

When executed, the Python script loaded the threat feed, read the input indicators, compared them, and generated a results file called `results.csv`. Indicators found in the threat feed were marked as known malicious or suspicious, while unknown indicators were marked as not found.

This module demonstrates the basic purpose of threat intelligence in cybersecurity: using known threat information to identify possible risks. It also shows how threat intelligence can support decision-making by helping analysts prioritise which indicators require further investigation.

Evidence Added

As evidence, I included screenshots of the input IOC list, the threat feed, the Python implementation, the terminal output, and the generated CSV report. These screenshots show both the design and implementation of the threat intelligence module. For full code pls go to my GitHub link:

https://github.com/faiyazkabr3/cyber_portfolio/blob/main/part_2/B25/EVIDENCE/threat_checker.py

```

1  import csv
2
3  def load_threat_feed(filename):
4      threat_data = {}
5
6      with open(filename, "r") as file:
7          reader = csv.DictReader(file)
8
9          for row in reader:
10             indicator = row["indicator"].strip()
11             threat_data[indicator] = {
12                 "type": row["type"],
13                 "threat_level": row["threat_level"],
14                 "description": row["description"]
15             }
16
17     return threat_data
18
19
20 def load_iocs(filename):
21     with open(filename, "r") as file:
22         return [line.strip() for line in file if line.strip()]
23
24
25 def check_iocs(iocs, threat_data):
26     results = []
27
28     for ioc in iocs:
29         if ioc in threat_data:
30             results.append({
31                 "indicator": ioc,
32                 "status": "Known malicious/suspicious",

```



threat_intel_module



iocs_to_check.txt



results.csv



threat_checker.py



threat_feed.csv

```

faiyazkahir@MacBook-Pro cyber_quiz_2 % cd threat_intel_module
[python3 threat_checker.py
Threat intelligence check completed.
Results saved to results.csv
185.220.101.1 - Not found - low
8.8.8.8 - Not found - low
malicious-example.com - Not found - low
google.com - Not found - low
bad-download-site.net - Not found - low
faiyazkahir@MacBook-Pro threat_intel_module %

```

	A	B	C	D	E	F	G	H	I	J	K
1	indicator	status	type	threat_level	description						
2	185.220.101.	Not found	unknown	low	No match found in local threat feed						
3	8.8.8.8	Not found	unknown	low	No match found in local threat feed						
4	malicious-ex	Not found	unknown	low	No match found in local threat feed						
5	google.com	Not found	unknown	low	No match found in local threat feed						
6	bad-downloa	Not found	unknown	low	No match found in local threat feed						
7											
8											
9											
10											
11											

B26. Help another student in this unit struggling to understand/learn a cybersecurity concept.

For this task, I helped another student in the unit, **Casey Zhu(24722864)**, who was struggling to understand the concept of an **Intrusion Detection System (IDS)**. I explained the topic from the beginning by first defining an IDS as a security tool that monitors network traffic or system activity to identify suspicious behaviour. I then explained that an IDS does not usually stop an attack by itself, but it helps detect possible attacks and creates alerts so that a person or security team can respond.

To make the concept easier to understand, I used a simple analogy: a firewall is like a locked door, while an IDS is like a security camera that watches for suspicious activity. I also explained the difference between **detection** and **response**, using the example of Suricata detecting suspicious network traffic and a firewall such as UFW being used afterward to block the source. This matched the IDS simulation I had worked on, where suspicious traffic was generated, Suricata produced IDS alerts, and UFW was used as the response control. The project report describes this defensive chain as “attack trigger, IDS detection, response action, and final proof” that the suspicious device could no longer reach critical services.

After the explanation, I showed **Casey Zhu** a simple example of how IDS alerts work, including how suspicious activity such as ping probing, HTTP access, SSH attempts, or SMB attempts can be detected using custom Suricata rules. In my IDS simulation, Suricata was used with custom rules to generate alerts for these types of activity. This helped **Casey Zhu** understand that an IDS is useful because it gives visibility into suspicious behaviour instead of only relying on prevention tools.

2.4 Detect and Respond to Attacks

Boba Tea Store Security Enhancement — IDS simulation, alerting, and containment

Kali attacker / delivery tablet

Ubuntu POS • Suricata IDS

UFW containment

```
ected [**] [Classification: (null)] [Priority: 3] [ICMP] 192.168.64.218 -> 192.168.64.3:0
35/01/2026-14:28:28.635820 [**] [1:1001001:1] BOBA ALERT: ICMP ping/probing detected [**] [Classification: (null)] [Priority: 3] [ICMP] 192.168.64.3:0 -> 192.168.64.218
35/01/2026-14:28:54.891896 [**] [1:1001001:1] BOBA ALERT: ICMP ping/probing detected [**] [Classification: (null)] [Priority: 3] [IPv6-ICMP] fe80:0000:0000:0000:0000:0000:0000:0000 -> fe02:0000:0000:0000:0000:0000:0000:0000:113a:0
35/01/2026-14:28:54.892209 [**] [1:1001003:1] BOBA ALERT: SSH access attempt to critical POS host [**] [Classification: (null)] [Priority: 3] [TCP] 192.168.64.1:55868 -> 192.168.64.3:22
35/01/2026-14:28:54.892209 [**] [1:1001004:1] BOBA ALERT: SMB access attempt to critical POS host [**] [Classification: (null)] [Priority: 3] [TCP] 192.168.64.1:35819 -> 192.168.64.3:445
35/01/2026-14:28:54.829585 [**] [1:1001003:1] BOBA ALERT: SSH access attempt to critical POS host [**] [Classification: (null)] [Priority: 3] [TCP] 192.168.64.1:55878 -> 192.168.64.3:22
35/01/2026-14:29:30.174286 [**] [1:1001002:1] BOBA ALERT: HTTP access to fake POS web server [**] [Classification: (null)] [Priority: 3] [TCP] 192.168.64.2:59516 -> 192.168.64.3:80
35/01/2026-14:30:04.540727 [**] [1:1001001:1] BOBA ALERT: ICMP ping/probing detected [**] [Classification: (null)] [Priority: 3] [IPv6-ICMP] fe80:0000:0000:0000:0000:0000:0000:0000:113a:0 -> fe02:0000:0000:0000:0000:0000:0000:0000:113a:0
35/01/2026-14:30:04.540727 [**] [1:1001001:1] BOBA ALERT: ICMP ping/probing detected [**] [Classification: (null)] [Priority: 3] [ICMP] 192.168.64.218 -> 192.168.64.3:0
```


Real IDS alert evidence from Suricata fast.log

Prepared by Faiyaz

Local lab: Kali Linux 192.168.64.2 → Ubuntu POS 192.168.64.3

Practical IDS deployment in the store

Keep the focus on detection and response, not just segmentation




Recommended placement options

- Simple: IDS/IPS on pfSense or OPNsense firewall.
- Stronger: passive Suricata sensor on a mirrored switch port.
- Confirm impact using POS host logs and firewall logs.

Operational principle

- Start detection-only to avoid breaking sales traffic.
- Tune alerts around POS, NVR, IoT and delivery devices.
- Enable IPS blocking only for high-confidence events.



B27. Apply a learned concept in this unit to a real-world application/problem/environment.

Learned Concept Applied: Threat Modelling

For this task, I applied the cybersecurity concept of **threat modelling** to a real-world environment that I use regularly: my **personal student laptop and online study accounts**. This includes my laptop, university account, email, GitHub account, cloud storage, browser sessions, and assignment files.

Threat modelling is useful because it helps identify what needs to be protected, what can go wrong, and what security controls should be applied. In the lecture, threat modelling is described as a process used to identify threats, vulnerabilities, missing safeguards, and prioritise countermeasures. The lecture also explains threat modelling using four key questions: **What are we working on? What can go wrong? What are we going to do about it? Did we do a good enough job?**

Real-World Environment

The real-world environment I analysed was my student study setup. This environment is important because it contains academic work, personal information, login credentials, project code, and assignment files. If this environment were compromised, it could lead to loss of coursework, account takeover, privacy issues, or exposure of sensitive project files.

Important Assets

Asset	Why It Matters
Laptop	Stores assignments, code, browser sessions, and personal files
University account	Provides access to LMS, emails, grades, and submissions
Email account	Used for password resets and important communications
GitHub account	Stores project code and collaboration history
Cloud storage	Stores backups and assignment files
Passwords/session cookies	Could allow account takeover if stolen
Project secrets	.env files or keys could expose private application data

Threat Modelling Tree

Main attack goal: Compromise my student study environment.

- |
- +-- 1. Gain access to my accounts
 - | +-- Steal university login (phishing, fake login page, reused password)
 - | +-- Gain access to email account (password reuse, credential stuffing, no MFA)
 - | +-- Access GitHub account (weak password, stolen session, no MFA)
- |
- +-- 2. Infect laptop with malware
 - | +-- Download unsafe software (cracked software, fake update, unofficial website)
 - | +-- Open malicious attachment (fake assignment file, scam PDF/ZIP)
 - | +-- Plug in unsafe USB or infected external drive
- |
- +-- 3. Steal or lose academic files
 - | +-- Malware/ransomware corrupts files
 - | +-- Laptop is lost or stolen
 - | +-- No cloud backup available
- |
- +-- 4. Leak sensitive information
 - | +-- Public Wi-Fi misuse or unsafe browsing
 - | +-- Commit .env/project secrets to GitHub
 - | +-- Session hijacking on shared/untrusted device

Threats Identified

Using the question “**What can go wrong?**”, I identified several realistic threats. The lecture explains that this question helps identify threats by considering intended users, unintended users, and unintended uses of a system.


Threat	Possible Impact
Phishing email stealing my university login	Unauthorised access to LMS, email, and submissions
Malware from unsafe downloads	File loss, credential theft, or system compromise
Weak or reused passwords	Easier account takeover
No backup of assignment files	Loss of coursework if laptop fails
Public Wi-Fi misuse	Exposure to unsafe network activity
Accidentally committing .env files to GitHub	Leakage of project secrets or keys
Stolen session cookies	Account access without needing the password

Mitigations Applied

The lecture explains that responses to threats can include **mitigating**, **eliminating**, **transferring**, or **accepting** threats. Mitigation means making it harder for someone to take advantage of a threat.

Threat	Mitigation
Phishing	Check sender address, avoid suspicious links, and use MFA
Malware	Download only from official websites and keep antivirus enabled
Weak passwords	Use long, unique passwords and avoid password reuse
File loss	Back up important files to cloud storage
Public Wi-Fi risk	Avoid logging into sensitive accounts on untrusted networks
GitHub secret leakage	Use .gitignore and avoid committing .env files
Session hijacking	Log out from shared devices and avoid saving sessions on public computers

Evidence:



Faiyaz Kabir (faiyazkabr3)

Your personal account

Go to your personal profile

Public profile

Account

Appearance

Accessibility

Notifications

Access

Billing and licensing

Emails

Password and authentication

Sessions

SSH and GPG keys

Organizations

Enterprises

Moderation

Code, planning, and automation

Repositories

Codespaces

Models

Packages

Copilot

Pages

Saved replies

Security

Code security

Integrations

Sign in methods

Email

2 verified emails configured

Manage

Password

Configured

Change password

Passkeys

1 passkey configured

Add passkey

Google

Sign in with your Google account

Connect

Apple

Sign in with your Apple account

Connect

Two-factor authentication

Two-factor authentication adds an additional layer of security to your account by requiring more than just a password to sign in. [Learn more about two-factor authentication.](#)

Preferred 2FA method

Set your preferred method to use for two-factor authentication when signing into GitHub.

Authenticator app

Two-factor methods

Authenticator app

Configured

Edit

Use an authentication app or browser extension to get two-factor authentication codes when prompted.

B28. Produce a cyber safety flyer for (choose 1): elders, high school students, CEOs, Uni students.




B29. Find a CVE in this year and fix it using three different generative AI systems (e.g., ChatGPT, Gemini), comparing the consistency.

Introduction

Cybersecurity vulnerabilities (CVEs) are weaknesses in software that attackers can exploit. In this task, I selected a recent CVE and used three generative AI systems to fix a similar vulnerability in code. I then compared the results to see how consistent the AI systems were.

Selected CVE

The CVE I selected is **CVE-2026-34621**.



This vulnerability affects Adobe Acrobat Reader and is caused by **improper handling of object properties (Prototype Pollution)**. If exploited, it could allow attackers to execute malicious code when a user opens a harmful file.

This type of issue is related to poor input validation and unsafe handling of data.

Vulnerable Code Example

To simulate this vulnerability, I created a simple JavaScript example that allows unsafe object modification:

```
1  function merge(target, source) {
2      for (let key in source) {
3          target[key] = source[key];
4      }
5      return target;
6  }
7
8  // Simulated malicious user input
9  let userInput = JSON.parse('{"__proto__": {"isAdmin": true}}');
10
11 // Target object
12 let obj = {};
13
14 // Merge user input into object
15 merge(obj, userInput);
16
17 // Check result
18 console.log(obj.isAdmin); // true (unexpected)
```

This code is vulnerable because it allows modification of the object prototype.

AI Systems Used

I used three AI tools:

- ChatGPT
- Google Gemini
- Microsoft Copilot

I gave all three the same prompt:

“Fix this JavaScript code to prevent prototype pollution and explain the solution.”

AI Fixes

1. ChatGPT Fix

ChatGPT suggested filtering dangerous keys like `__proto__`, `constructor`, and `prototype`

```
1 function merge(target, source) {
2   for (let key in source) {
3     // Skip dangerous keys
4     if (key === "__proto__" || key === "constructor" || key === "prototype") {
5       continue;
6     }
7
8     target[key] = source[key];
9   }
10
11   return target;
12 }
```

2. Gemini Fix

Gemini suggested using safer object handling and validating input before merging

```
1 function merge(target, source) {
2   Object.keys(source).forEach(function(key) {
3     // Only allow safe keys
4     if (!["__proto__", "constructor", "prototype"].includes(key)) {
5       target[key] = source[key];
6     }
7   });
8
9   return target;
10 }
```

3. Copilot Fix

Copilot suggested both filtering keys and using `hasOwnProperty` for safer access.

```
1 function merge(target, source) {
2   for (let key in source) {
3     // Check if property belongs to object
4     if (Object.prototype.hasOwnProperty.call(source, key)) {
5
6       // Block dangerous keys
7       if (key !== "__proto__" && key !== "constructor" && key !== "prototype") {
8         target[key] = source[key];
9       }
10    }
11  }
12
13  return target;
14 }
```

Comparison of AI Results:

Criteria	ChatGPT	Gemini	Copilot
Prevented prototype pollution	Yes	Yes	Yes

Filtered dangerous keys	Yes	Yes	Yes
Additional safety checks	Moderate	Moderate	Strong
Explanation clarity	High	Medium	Medium

B30. Generate an AI-created image, applying an imperceptible watermark on it and then perform an image-to-image regeneration or editing process to make sure the watermark is detectable---the watermark survives.

For this task, I generated an AI-created image and applied an imperceptible watermark to it. The aim was to test whether the watermark could still be detected after the image was edited or regenerated.

First, I created an AI-generated image and saved it as the original image. I then used a Python watermarking script to embed a hidden watermark into the image. The watermark used for this test was “C”. The watermark was designed to be invisible to the human eye, meaning the image still looked normal after the watermark was added.

After embedding the watermark, I performed an image editing process on the watermarked image. This was done to test whether the watermark could survive image modification. The edited version was then checked using a watermark detection script.

The detection result showed:

Decoded watermark from edited image: C

This confirmed that the watermark was still detectable after the image had been edited.

Therefore, the test was successful because the imperceptible watermark survived the image editing process.

Evaluation

This activity showed that invisible watermarking can be used to embed hidden ownership or verification information inside AI-generated images. The watermark was not visually obvious, but it could still be detected using the correct decoding script. This is useful because watermarking can help identify AI-generated content or prove that an image has been previously marked.

However, the test also showed that watermark survival depends on the strength of the image modification. A light edit preserved the watermark, but heavier image-to-image regeneration, compression, cropping, or major visual changes may damage or remove it. Therefore, invisible watermarking is useful, but it is not always guaranteed to survive every type of image transformation.

Evidence Added


For evidence, I included:

- the original AI-generated image
- the watermarked image
- the edited/regenerated image
- the Python watermark embedding script
- the Python watermark detection script
- a terminal screenshot showing:

Decoded watermark from edited image: C

all evidence is here:

https://github.com/faiyazkabr3/cyber_portfolio/tree/main/part_2/B30/evidence



NAME	FILE SIZE	TYPE	DATE
original.png	7/05/2026 3:58 PM	PNG File	2,203 KB
watermarked	7/05/2026 4:12 PM	PNG File	1,902 KB

```

7 PS C:\Users\faiyaz> python detect_watermark.py
8 Decoded watermark from edited image: C
PS C:\Users\faiyaz>

```

```

embed_watermark.py 2 X detect_watermark.py 2
C:\Users> Users > faiyaz > embed_watermark.py > ...
1 import cv2
2 from imwatermark import WatermarkEncoder, WatermarkDecoder
3
4 # watermark text
5 wm_text = "C"
6
7 # load original image
8 img = cv2.imread(r"C:\Users\faiyaz\original.png.png")
9
10 # create encoder
11 encoder = WatermarkEncoder()
12 encoder.set_watermark("bytes", wm_text.encode("utf-8"))
13
14 # embed invisible watermark
15 watermarked = encoder.encode(img, "dwtDct")
16
17 # save result
18 cv2.imwrite(r"C:\Users\faiyaz\watermarked.png", watermarked)
19
20 print("Watermark embedded and saved as watermarked.png")

```

References (APA 7th Style)

- Australian Cyber Security Centre. (2023). *Essential Eight maturity model*.
<https://www.cyber.gov.au/resources-business-and-government/essential-cyber-security/essential-eight>
- Australian Cyber Security Centre. (2026). *Alerts and advisories*.
<https://www.cyber.gov.au/about-us/view-all-content/alerts-and-advisories>
- Cloudflare, Inc. (2023). *Web application firewall (WAF)*.
<https://www.cloudflare.com>
- National Institute of Standards and Technology. (2001). *Announcing the Advanced Encryption Standard (AES) (FIPS PUB 197)*.
<https://doi.org/10.6028/NIST.FIPS.197>
- Rescorla, E. (2018). *The Transport Layer Security (TLS) protocol version 1.3 (RFC 8446)*.
<https://www.rfc-editor.org/rfc/rfc8446>
- Whitman, M. E., & Mattord, H. J. (2022). *Principles of information security* (7th ed.). Cengage Learning.
- Australian Human Rights Commission. (n.d.). *Technology and human rights: Historical bias in AI systems*. Retrieved from <https://humanrights.gov.au/about-us/media-centre/media-releases/technology-and-human-rights/historical-bias-ai-systems>
- Zhang, Y., Tan, L., & O'Halloran, K. L. (2026). *Bias in GenAI: Exploring cultural variations between large language models*. *Linguistics and Education*, 92, 101515.
<https://doi.org/10.1016/j.linged.2026.101515>
- Cybersecurity and Infrastructure Security Agency. (n.d.-b). Require multifactor authentication. <https://www.cisa.gov/resources-tools/resources/require-multifactor-authentication>
- MITRE. (n.d.-a). About the D3FEND knowledge graph project.
<https://d3fend.mitre.org/about/>

MITRE. (n.d.-b). ATT&CK. <https://attack.mitre.org/>

MITRE. (n.d.-c). D3FEND matrix. <https://d3fend.mitre.org/>

National Institute of Standards and Technology. (2020). Zero trust architecture (SP 800-207). <https://csrc.nist.gov/publications/detail/sp/800-207/final>